

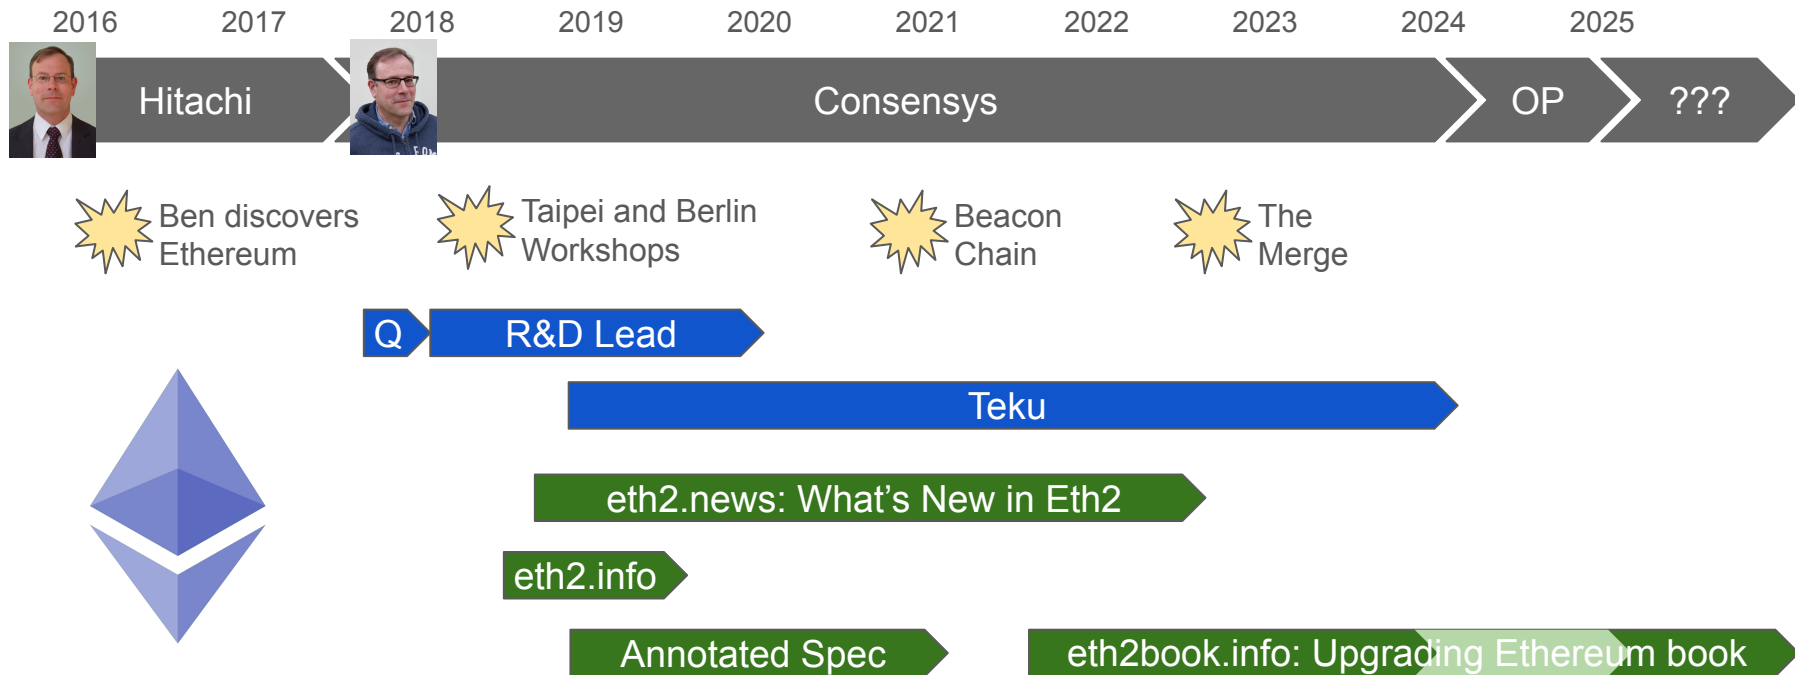
# Fork Choice

Ethereum Protocol Study Group

2025-03-10



# Ben's adventures in Ethereum



# Agenda

1. Introduction to Fork Choice
2. LMD GHOST
3. Casper FFG
4. Gasper
5. AMA


# Introduction

# Recap: Safety and Liveness in distributed systems

- Safety: consistency / finalisation
  - “Nothing bad ever happens”
- Liveness: availability
  - “Something good eventually happens”
- CAP Theorem
  - You can’t guarantee both! (On a real network)
- Ethereum prioritises liveness
  - LMD GHOST
- With “best efforts” safety / finality
  - Casper FFG - a “finality gadget”
- In combination: GASPER

# The fork choice rule


- Liveness prioritising networks are “forkful”
  - They (sometimes) build trees rather than chains.
- A symptom of the lack of safety: different nodes have different views of the network.



# The fork choice rule

The “fork choice rule” is the method by which the network converges on a single branch, and therefore a single, shared history. Block tree → Blockchain.

Heuristic: build on the branch that is most likely to be built on (least likely to be orphaned) by other nodes, based on whatever information we have received.



# Here be Dragons!

Historical issues with Ethereum's PoS fork choice:

- August 2019, [decoy flip-flop attack](#) (temporary finality delay)
- September 2019, [bouncing attack](#) (indefinite finality delay)
- November 2019, [finality liveness failure](#) (unable to finalise new checkpoints)
- October 2020, [balancing attack](#) (long forks, indefinite finality delay)
- July 2021, [inconsistent finalised and justified](#)
- November 2021, [non-atomic finalisation and justification](#)
- January 2022, [a new balancing attack](#)
- May 2022, [unrealised justification reorgs](#) (proposers can induce long reorgs)
- June 2022, [justification withholding attacks](#) (ditto)


[Reference](#)



# When fork choice struggles: Holesky / Pectra upgrade

“40 hours in the life of fork choice” [interop-Discord channel](#) (2025-02-24 22:00)

- Faulty clients
- P2P meltdown
- Slashing protection
- Low block production
- Snap sync
- Checkpoint sync
- Optimistic sync
- Client instability (oom)
- ...



h/t [smartprogrammer.eth](https://smartprogrammer.eth)

# Navigating the specs

```
class AttestationData(Container):  
    slot: Slot  
    index: CommitteeIndex  
    # LMD GHOST vote  
    beacon_block_root: Root  
    # FFG vote  
    source: Checkpoint  
    target: Checkpoint
```

Each validator broadcasts its view of the network via an attestation, once per epoch (32 slots, 6.4 minutes):

- Broadcast via P2P gossip
  - Packed into blocks by proposers
- 
- LMD GHOST largely depends on attestations received via gossip
    - Therefore handling is part of fork choice: *fork-choice.md*
  - Casper FFG solely depends on attestations in blocks
    - Therefore handling is part of block and epoch processing: *beacon-chain.md*

# The Store

Each node maintains its own view of the network in an object called [the Store](#).

```
@dataclass
class Store(object):
    time: uint64
    genesis_time: uint64
    justified_checkpoint: Checkpoint
    finalized_checkpoint: Checkpoint
    * unrealized_justified_checkpoint: Checkpoint
    * unrealized_finalized_checkpoint: Checkpoint
    * proposer_boost_root: Root
    * equivocating_indices: Set[ValidatorIndex]
    blocks: Dict[Root, BeaconBlock] = field(default_factory=dict)
    block_states: Dict[Root, BeaconState] = field(default_factory=dict)
    * block_timeliness: Dict[Root, boolean] = field(default_factory=dict)
    checkpoint_states: Dict[Checkpoint, BeaconState] = field(default_factory=dict)
    latest_messages: Dict[ValidatorIndex, LatestMessage] = field(default_factory=dict)
    * unrealized_justifications: Dict[Root, Checkpoint] = field(default_factory=dict)
```

\* things added later  
to fix issues

# Events

- The LMD GHOST fork choice is event driven.
- Handlers:
  - [on\\_tick\(\)](#)
    - Regularly updates the current time and does some light housekeeping.
  - [on\\_block\(\)](#)
    - Adds new blocks to the Store as the node receives them
  - [on\\_attestation\(\)](#)
    - Updates validators' latest messages, whether received in blocks or via P2P
  - [on\\_attester\\_slashing\(\)](#)
    - Ensures that slashed validators' votes are not counted, and so avoids some equivocation attacks.
- It is always ready to output a best head block via a call to [get\\_head\(\)](#)

LMD GHOST

# Naming

- LMD

- “Message Driven”
  - Relies only on attestations
  - Validators attest to what they believe to be the best head in the current slot
- “Latest”
  - Only the most recent attestation from each validator counts
  - IMD GHOST (immediate), FMD GHOST (fresh), RLMD GHOST (recent latest)...

- GHOST


- “Greedy Heaviest-Observed Sub-Tree” algorithm
- from a 2013 paper by Sompolinsky and Zohar about how to safely improve transaction throughput on Bitcoin

```
class AttestationData(Container):  
    slot: Slot  
    index: CommitteeIndex  
    # LMD GHOST vote  
    beacon_block_root: Root  
    # FFG vote  
    source: Checkpoint  
    target: Checkpoint
```

# LMD GHOST Overview

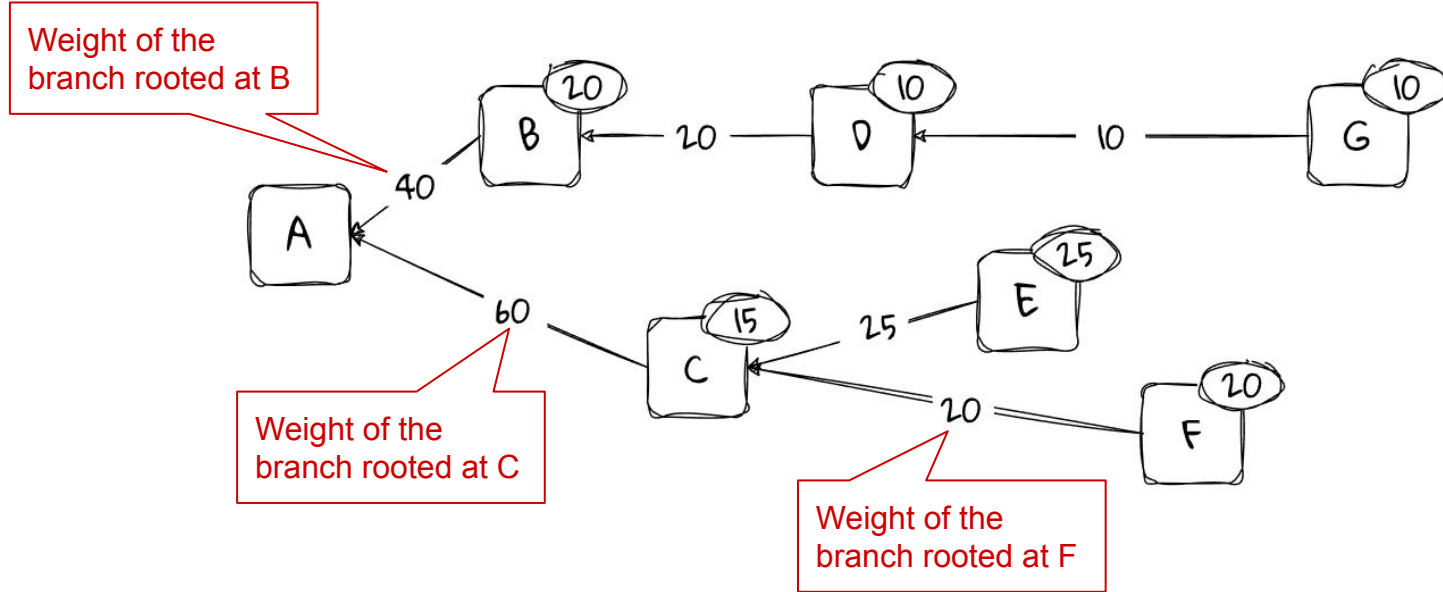
- Timescale:
  - Slot-based: 12s
- Goal:
  - Used by the block proposer to decide on which branch to build its block
  - Used by attesters to choose which branch to attest to => convergence
- Heuristic:
  - Which branch is least likely to be orphaned in future?
- Based on:
  - “Weighing” the votes received in attestations
  - Note that a maximum of only 3.125% of the votes are “fresh” (same-slot)
- Properties
  - Liveness: it will always output a viable head block on which to build ([proof under synchrony](#))
  - Safety: no useful guarantees (however, see the [confirmation rule](#) work.)

# LMD GHOST fork choice






# LMD GHOST fork choice



# LMD GHOST fork choice



# get\_head()

```
def get_head(store: Store) -> Root:
    # Get filtered block tree that only includes viable branches
    blocks = get_filtered_block_tree(store)
    # Execute the LMD-GHOST fork choice
    head = store.justified_checkpoint.root
    while True: (recursive function written iteratively)
        children = [
            root for root in blocks.keys()
            if blocks[root].parent_root == head
        ]
        if len(children) == 0:
            return head
        # Sort by latest attesting balance with ties broken lexicographically
        # Ties broken by favoring block with lexicographically higher root
        head = max(children, key=lambda root: (get_weight(store, root), root))
```

In pure LMD GHOST this would be the genesis block

Stop when we reach a branch tip

Follow the branch with the maximum weight

Weigh the branch rooted at block "root"

Phase 0  
[fork-choice.md](#)

[Implementations](#)

# get\_weight() (of a branch rooted at block “root”)


```
def get_weight(store: Store, root: Root) -> Gwei:
    state = store.checkpoint_states[store.justified_checkpoint]
    unslashed_and_active_indices = [
        i for i in get_active_validator_indices(state, get_current_epoch(state))
        if not state.validators[i].slashed
    ]
    attestation_score = Gwei(sum(
        state.validators[i].effective_balance for i in unslashed_and_active_indices
        if (i in store.latest_messages
            and i not in store.equivocating_indices
            and get_ancestor(store, store.latest_messages[i].root, store.blocks[root].slot) == root)
    ))
    if store.proposer_boost_root == Root():
        # Return only attestation score if ``proposer_boost_root`` is not set
        return attestation_score
```

All validators eligible to vote

Votes are weighted by EB

Filter on validators whose latest vote was for a block whose ancestor is root

# LMD GHOST fork choice



# Ignoring undesirable blocks

- Was it correctly signed by the expected proposer?
  - Actually checked upstream of fork choice (gossip layer?)
- Are all its ancestor blocks valid?
- Is the block's post-state hash correct?
  - That is, does the block respect the STF (State Transition Function)?

## New in Deneb:

- Is its data available?
  - Have we seen/can we get its attached blobs?

## Potential future:

- Is the block non-censoring?
- More sophisticated data availability (e.g. PeerDAS - see [Fulu](#))

# Ignoring undesirable blocks

```
def on_block(store: Store, signed_block: SignedBeaconBlock) -> None:
    """
    Run ``on_block`` upon receiving a new block.
    """
    block = signed_block.message
    # Parent block must be known
    assert block.parent_root in store.block_states
    # Blocks cannot be in the future. If they are, their consideration must be delayed until they are in the past.
    assert get_current_slot(store) >= block.slot

    # Check that block is later than the finalized epoch slot (optimization to reduce calls to get_ancestor)
    finalized_slot = compute_start_slot_at_epoch(store.finalized_checkpoint.epoch)
    assert block.slot > finalized_slot
    # Check block is a descendant of the finalized block at the checkpoint finalized slot
    finalized_checkpoint_block = get_checkpoint_block(
        store,
        block.parent_root,
        store.finalized_checkpoint.epoch,
    )
    assert store.finalized_checkpoint.root == finalized_checkpoint_block

    # [New in Deneb:EIP4844]
    # Check if blob data is available
    # If not, this block MAY be queued and subsequently considered when blob data becomes available
    # *Note*: Extraneous or invalid Blobs (in addition to the expected/referenced valid blobs)
    # received on the p2p network MUST NOT invalidate a block that is otherwise valid and available
    assert is_data_available(hash_tree_root(block), block.body.blob_kzg_commitments)
```

} Are all its ancestors viable?

Is not from the future

} Is not too old

Data is available


Deneb  
[fork-choice.md](#)

# Proposer Boost

- In 2020 a “balancing attack” was identified
  - A (very) small number of validators could fork the chain into two equally weighted branches and maintain this indefinitely.
- The defence against this is “Proposer Boost”
  - Blocks received in a timely way (within 4 seconds of slot start) are assigned a huge extra weight in the `get_weight()` calculation.
    - The extra weight almost certainly makes the timely block the head of the chain
  - This makes honest proposers almost sure to build on a timely block that they received
  - “A block that is timely should not expect to be reorged”



# Proposer Boost and fork choice



Why 40?

# Proposer Boost in code

## get\_weight()

Dummy value

```
if store.proposer_boost_root == Root():  
    # Return only attestation score if ``proposer_boost_root`` is not set  
    return attestation_score  
  
# Calculate proposer score if ``proposer_boost_root`` is set  
proposer_score = Gwei(0)  
# Boost is applied if ``root`` is an ancestor of ``proposer_boost_root``  
if get_ancestor(store, store.proposer_boost_root, store.blocks[root].slot) == root:  
    proposer_score = get_proposer_score(store)  
return attestation_score + proposer_score
```

```
def get_proposer_score(store: Store) -> Gwei:  
    justified_checkpoint_state = store.checkpoint_states[store.justified_checkpoint]  
    committee_weight = get_total_active_balance(justified_checkpoint_state) // SLOTS_PER_EPOCH  
    return (committee_weight * PROPOSER_SCORE_BOOST) // 100
```

40

Phase 0  
[fork-choice.md](#)

Add 40% of the total stake attesting in this slot to the weight of the branch with the timely block.

# Proposer Boost in code

## on\_block()

```
# Add block timeliness to the store
time_into_slot = (store.time - store.genesis_time) % SECONDS_PER_SLOT
is_before_attesting_interval = time_into_slot < SECONDS_PER_SLOT // INTERVALS_PER_SLOT
is_timely = get_current_slot(store) == block.slot and is_before_attesting_interval
store.block_timeliness[hash_tree_root(block)] = is_timely

# Add proposer score boost if the block is timely and not conflicting with an existing block
is_first_block = store.proposer_boost_root == Root()
if is_timely and is_first_block:
    store.proposer_boost_root = hash_tree_root(block)
```

A timely block is received  
within 4s of the slot start

Record the timely block for use  
in get\_weight()

Phase 0  
[fork-choice.md](#)

## on\_tick\_per\_slot()

```
# If this is a new slot, reset store.proposer_boost_root
if current_slot > previous_slot:
    store.proposer_boost_root = Root()
```

Dummy value

# Forking out late blocks with Proposer Boost

- Proposer Boost solves the balancing attack, but also gives us the opportunity to strongly discourage late block publication.
  - Honest proposers should publish a block by 4s to allow time for attestations to be made and to be collected by the network.
  - However, the spec allows a block received as late as 12s to be built upon by the next proposer.
  - Block builders (MEV searchers) took advantage of this to publish later and later with low risk of being orphaned.
  - This destabilises the network as such late blocks gather few attestations, and it penalises timely attesters (who attest to an empty slot)
- Proposer Boost allows an honest proposer to build on the *parent* of a late block, its block becoming the head (since it is timely)
  - Thus the late block is orphaned / re-orged out of the chain
  - See [get\\_proposer\\_head\(\)](#)

# Slashing in LMD GHOST

Slashing solves the “nothing at stake” problem with proof of stake.

Equivocation is costless:

- It's costless for proposers to publish multiple blocks (unlike in PoW)
  - E.g. why not build on every branch?
- It's costless for validators to vote for everything they see
  - A rich source of fork choice attacks.

Solution:

- Slash equivocating proposers and equivocating attestors

January 2014: [Slasher: A Punitive Proof-of-Stake Algorithm](#)

# Yay! We have ourselves a consensus mechanism 🎉

- Timescale:
  - Slot-based: 12s
- Goal:
  - Used by the block proposer to decide on which branch to build its block
  - Used by attesters to choose which branch to attest to => convergence
- Heuristic:
  - Which branch is least likely to be orphaned in future?
- Based on:
  - “Weighing” the votes received in attestations
  - Note that a maximum of only 3.125% of the votes are “fresh” (same-slot)
- Properties
  - Liveness: it will always output a viable head block on which to build ([proof under synchrony](#))
  - Safety: no useful guarantees (however, see the [confirmation rule](#) work.)

# But we can do better 🤔

Under proof of stake:

- 1) We have a well-defined set of participants
  - as in classical consensus, unlike PoW
- 2) We can apply economic penalties for bad behaviour
  - slashing

How about we use these to add some stronger safety properties to our chain?

# Casper FFG



# Naming



- Casper
  - Kind of influenced by Vlad Zamfir's Casper protocol that uses GHOST
    - see Casper the Friendly Ghost...
  - But Casper FFG has very little similarity to Zamfir's protocols and doesn't use GHOST 🙅
- FFG
  - The "Friendly Finality Gadget"
    - A "gadget" in Vitalik-speak is a self-contained enhancement to another process
    - So, Casper FFG is a gadget for adding finality to an existing blockchain consensus mechanism
    - Was originally planned to add finality to PoW as part of the PoS transition (EIP-1011)


[Original Casper FFG Paper](#)

# Casper FFG Overview

- Timescale:
  - Epoch-based: 32 slots / 6.4 minutes / 384 s
- Goal:
  - Confer “finality” on the chain: a checkpoint that will never be reverted (except at enormous cost)
- Heuristic:
  - Two-phase commit based on agreement among validators having at least  $\frac{2}{3}$  of the stake
- Based on:
  - “Weighing” the source and target votes received in attestations contained in blocks
- Properties
  - [Plausible Liveness](#): cannot get into a stuck state that is unable to finalise anything
  - [Accountable Safety](#): finalising conflicting checkpoints comes at enormous cost
    - aka Economic Finality

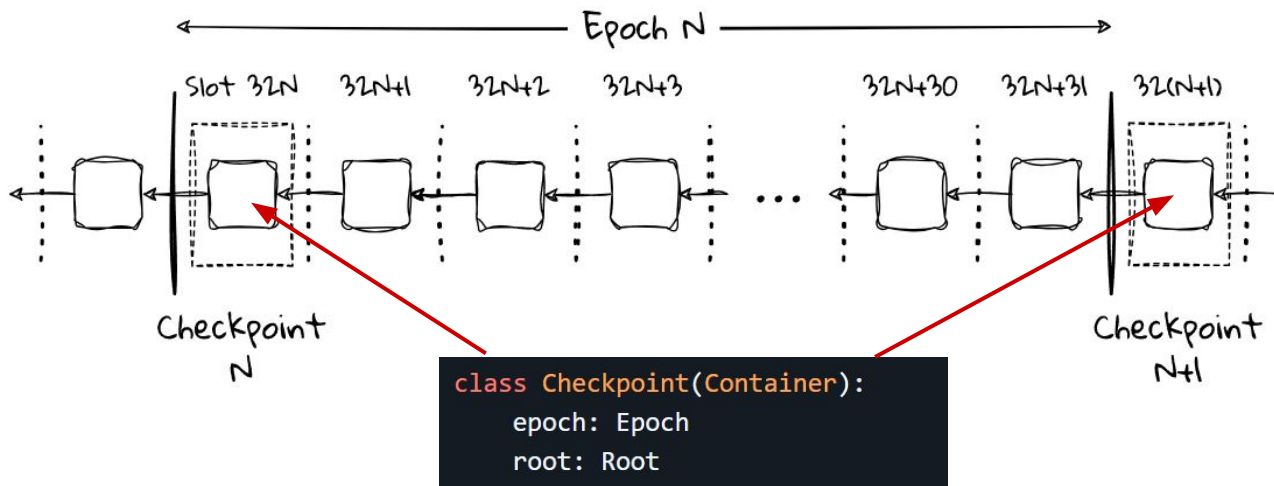
# Checkpoints

- Casper FFG relies on seeing votes from the whole validator set.
- This is too expensive to do every slot, so we accumulate votes across an epoch.
- 1/32 of the validator set votes at each slot.
- Accounting is done at each epoch end transition.



# Checkpoints


- A checkpoint is the block at the first slot of the Epoch
  - Or the previous block if that slot is empty



# Checkpoints

- Validators vote for a source and a target checkpoint
  - Source: the highest justified checkpoint I know of.
  - Target: the checkpoint I see in the current epoch.


```
class AttestationData(Container):  
    slot: Slot  
    index: CommitteeIndex  
    # LMD GHOST vote  
    beacon_block_root: Root  
    # FFG vote  
    source: Checkpoint  
    target: Checkpoint
```



# Supermajority links


- A supermajority link occurs when  $> \frac{2}{3}$  of validators (by stake) have made the same source  $\rightarrow$  target vote

```
class AttestationData(Container):  
    slot: Slot  
    index: CommitteeIndex  
    # LMD GHOST vote  
    beacon_block_root: Root  
    # FFG vote  
    source: Checkpoint  
    target: Checkpoint
```



# Two-phase commit: justification


- My target vote says “My preferred branch has current checkpoint X”.
- If I hear that  $\frac{2}{3}$  of validators agree with me, I locally mark X *justified*.
  - I’ve seen a supermajority link with X as target



# Two-phase commit: finalisation


- My source vote says “I heard that  $>\frac{2}{3}$  of you like X so I justified it”
- If I hear that  $\frac{2}{3}$  of validators also said that, I locally mark X as *finalised*.
  - I’ve seen a supermajority link with X as source
  - We know that  $\frac{2}{3}$  of validators have justified X and will therefore never revert it

## Target votes



Round 1  
Justification

## Source votes




Round 2  
Finalisation




# Casper Commandment: no double vote

A validator must not publish distinct votes  
 $s_1 \rightarrow t_1$  and  $s_2 \rightarrow t_2$  such that  $h(t_1) = h(t_2)$

That is, a validator must make at most one vote for any target epoch.



No!




No!


# Casper Commandment: no surround vote

A validator must not publish distinct votes  $s_1 \rightarrow t_1$  and  $s_2 \rightarrow t_2$  such that  $h(s_1) < h(s_2) < h(t_1) < h(t_2)$

That is, a validator must not make a vote such that its link either surrounds, or is surrounded by, a previous link it voted for.



No!




No!

(See Holesky)

# Accountable Safety and Plausible Liveness

- The proofs of [Accountable Safety](#) and [Plausible Liveness](#) rely on these two commandments.
- And one further rule that we only finalise a checkpoint that is the direct child of a justified checkpoint



Except, see [k-finality](#)

# Slashing in Casper FFG


Delivers “economic finality”: we slash validators that break a commandment

- If  $< \frac{1}{3}$  of validators is adversarial, safety is guaranteed (as with PBFT)
- If  $> \frac{1}{3}$  of validators acts so as to finalise conflicting checkpoints:
  - At least  $\frac{1}{3}$  of validators must have broken a commandment.
  - We can detect this and [prove it onchain!](#)
  - The aggregate cost to the attacker would be at least  $\frac{1}{3}$  of the total staked Ether
  - Thus, we have “economic finality” - a quantifiable cost to messing with finality, even if the attacker has more than  $\frac{1}{3}$  of the validators


This is a significant feature: classical consensus makes no guarantees if  $> \frac{1}{3}$  of replicas are adversarial / faulty / Byzantine. BFT algorithms will also halt if  $> \frac{1}{3}$  of replicas are offline, whereas Casper FFG continues trying to justify checkpoints.

GASPER

# Applying the FFG to LMD GHOST



# Modified LMD GHOST fork choice



Apply Casper FFG's fork choice: "follow the chain containing the justified checkpoint of the greatest height".

# get\_head()

```
def get_head(store: Store) -> Root:
    # Get filtered block tree that only includes viable branches
    blocks = get_filtered_block_tree(store)
    # Execute the LMD-GHOST fork choice
    head = store.justified_checkpoint.root
    while True:
        children = [
            root for root in blocks.keys()
            if blocks[root].parent_root == head
        ]
        if len(children) == 0:
            return head
        # Sort by latest attesting balance with ties broken lexicographically
        # Ties broken by favoring block with lexicographically higher root
        head = max(children, key=lambda root: (get_weight(store, root), root))
```

Apply Casper FFG fork choice

Follow the branch with  
the maximum weight

Weigh the branch  
rooted at block "root"

Phase 0

[fork-choice.md](#)



The world were that the end of the story...



# Gaspar in reality



**Interaction bugs.** The “interface” between Casper FFG finalization and LMD GHOST fork choice is a source of significant complexity, leading to a number of attacks that have required fairly complicated patches to fix, with more weaknesses being regularly discovered.

[Vitalik](#)

How LMD GHOST and Casper FFG are actually bolted together

# Issue one: block tree filtering

```
def get_head(store: Store) -> Root:
    # Get filtered block tree that only includes viable branches
    blocks = get_filtered_block_tree(store)
    # Execute the LMD-GHOST fork choice
    head = store.justified_checkpoint.root
    while True:
        children = [
            root for root in blocks.keys()
            if blocks[root].parent_root == head
        ]
        if len(children) == 0:
            return head
        # Sort by latest attesting balance with ties broken lexicographically
        # Ties broken by favoring block with lexicographically higher root
        head = max(children, key=lambda root: (get_weight(store, root), root))
```

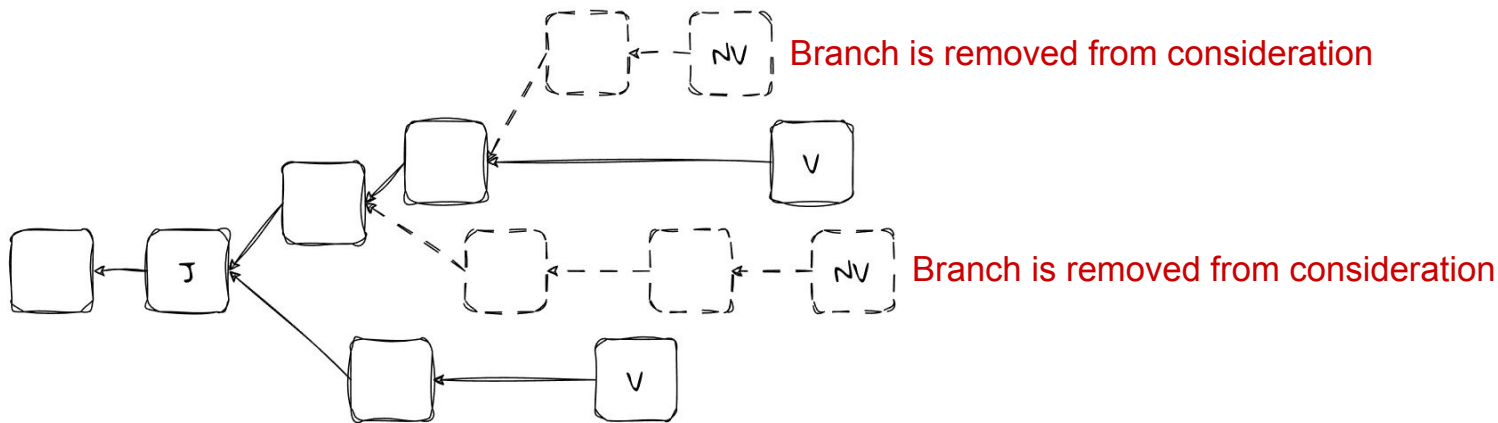
WTF is this???

Phase 0

[fork-choice.md](#)

## Issue one: block tree filtering

- filter\_block\_tree() removes from fork choice consideration any branches with “non-viable heads”.

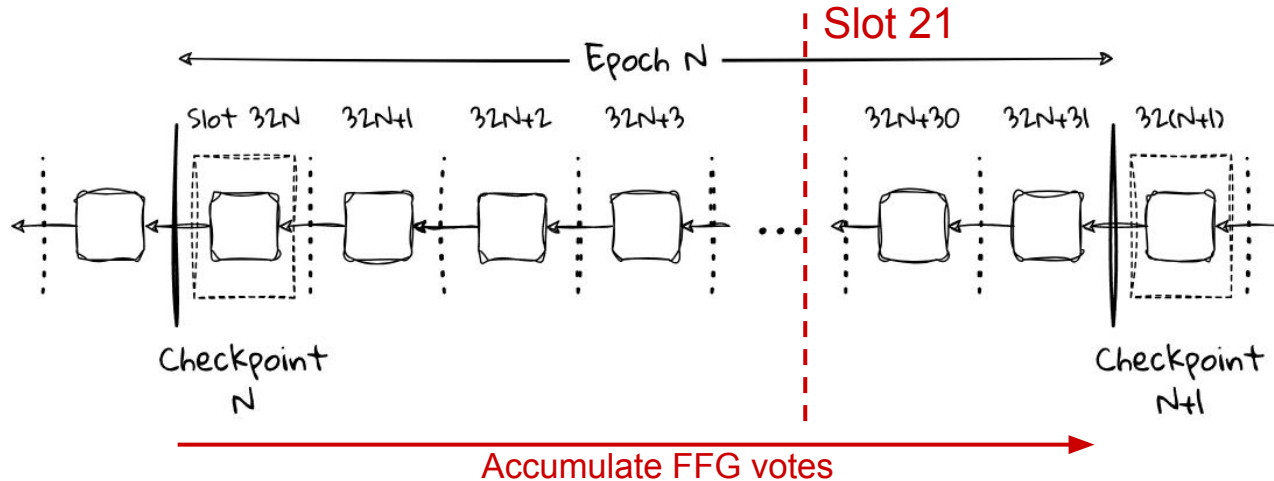


# Issue one: block tree filtering

- Non-viable means that the block's associated state does not agree with my Store about the current justified and finalised checkpoints.
- Resolves a potential [finalisation deadlock issue](#).
  - Plausible Liveness failure.
- Originates from a conflict between LMD GHOST's fork choice and Casper FFG's fork choice:
  - Casper's rule is "start from the highest justified checkpoint"
  - How to handle this in LMD GHOST when different blocks have different views of the highest justified checkpoint?

## Issue two: unrealised justification and finalisation

- Justification and finalisation accounting is done at the end of each epoch.
- However, by slot 21 ( $\frac{2}{3}$  of the way through the epoch) I might have seen enough votes to form a supermajority link: *unrealised* justification



# Issue two: unrealised justification and finalisation

- Unrealised justification reorg attack
  - allows the proposer of the first block of an epoch to fork out up to nine blocks from the end of the previous epoch.
- Justification withholding attack
  - adversary can reorg arbitrary numbers of blocks at the start of an epoch.
- Main issue:
  - Justification is updated in the beacon state only at the end of an epoch
  - Filter block tree was too aggressive because it did not account for “unrealised justification”
    - Blocks that would match my checkpoints if they had gone through epoch processing
    - Marked as non-viable, while actually being viable, and being removed from fork choice
      - Leads to re-orgs.

Unrealised Justification



The world when we have single slot finality

